

System Requirements

```
In [ ]: import sys
print("Python version: {}".format(sys.version))
import pandas as pd
print("pandas version: {}".format(pd.__version__))
import matplotlib
import matplotlib.pyplot as plt
print("matplotlib version: {}".format(matplotlib.__version__))
import numpy as np
print("NumPy version: {}".format(np.__version__))
import scipy as sp
print("SciPy version: {}".format(sp.__version__))
# import IPython
# print("IPython version: {}".format(IPython.__version__))
import sklearn
print("scikit-learn version: {}".format(sklearn.__version__))
import mglearn
print("mglearn version: {}".format(mglearn.__version__))
%matplotlib inline
```

Python version: 3.12.4 | packaged by Anaconda, Inc. | (main, Jun 18 2024, 15:03:56) [MSC v.1929 64 bit (AMD64)]
pandas version: 2.2.2
matplotlib version: 3.8.4
NumPy version: 1.26.4
SciPy version: 1.13.1
scikit-learn version: 1.4.2
mglearn version: 0.2.0

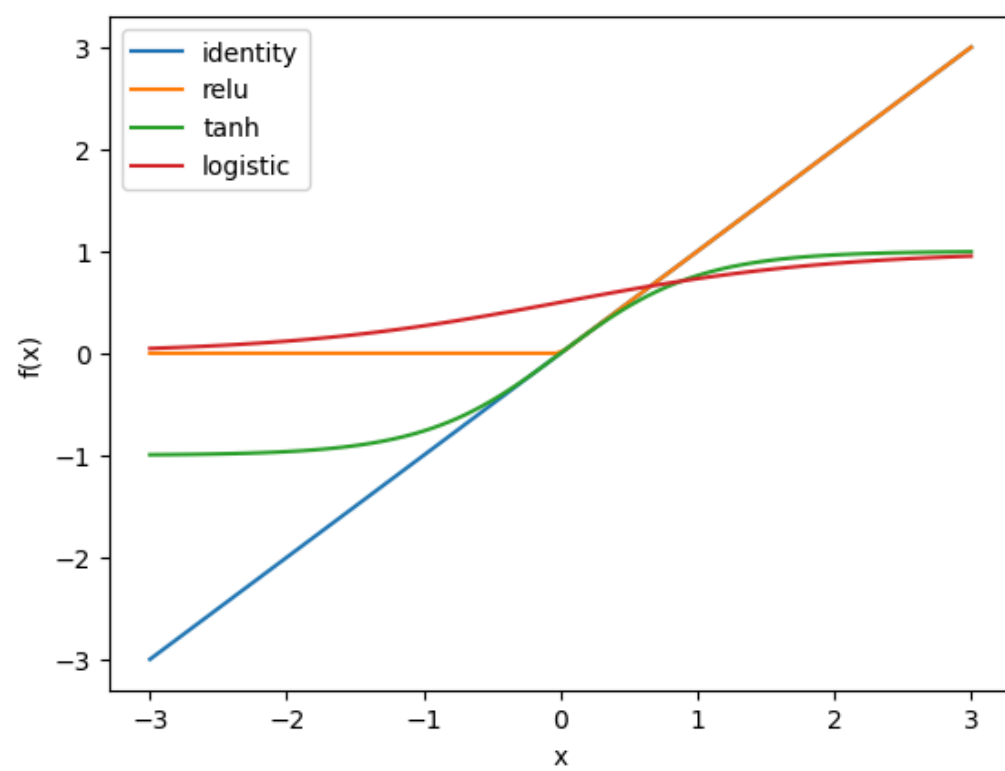
Introduction

Activation function


- Activation functions of a node defines the output of that node given an input or set of inputs e.g. MLPClassifier supports 'relu', 'identity', 'logistic' (=sigmoid), 'tanh'.
- This activation functions makes deep learning different from simply doing iterative LR on each layer, e.g, Tanh contains the output to between -1 and 1.
- Discover which activation function is best by trial and error

```
In [ ]: line = np.linspace(-3, 3, 100)
plt.plot(line, line, label="identity")
plt.plot(line, np.maximum(line, 0), label="relu")
plt.plot(line, np.tanh(line), label="tanh")
plt.plot(line, 1 / (1 + np.exp(-line)), label="logistic")
plt.legend(loc="best")
plt.xlabel("x")
plt.ylabel("f(x)")
```

Out[]: Text(0, 0.5, 'f(x)')




Perceptron




Gradient Descent

1. Initialize the weights with random values.
2. Choose a learning rate between 0 to 1: Low learning rate is slower but more accurate than high learning rate
3. Till the error is almost constant:
 - 3.1 calculate change in weight Δw
 - 3.2 update the weight

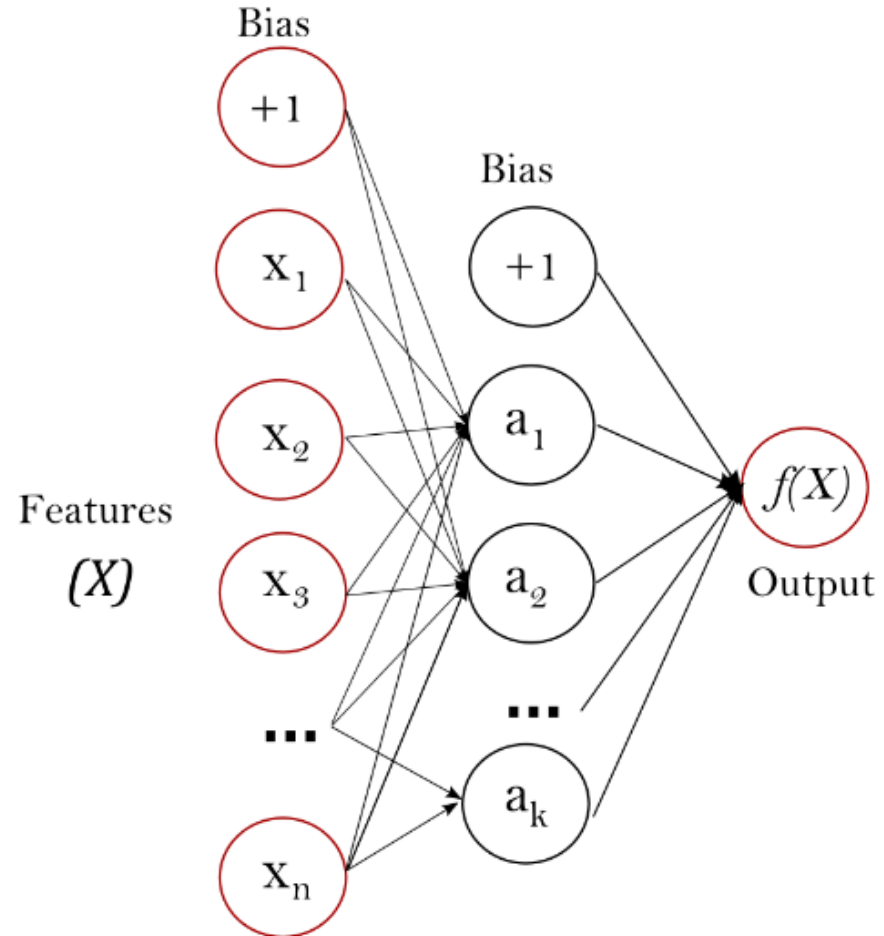
Gradient Step ($\|g\|=1.281474e+01$)



Newton Step ($\|g\|=1.232242e+01$)




Neural Network: Multi-layer Perceptron



Backpropagation

Visualization

ANN training



Initialize all weights to small random numbers.
Until satisfied, Do

- For each training example, Do
 1. Input the training example to the network and compute the network outputs
 2. For each output unit k

$$\delta_k \leftarrow o_k(1 - o_k)(t_k - o_k)$$

3. For each hidden unit h

$$\delta_h \leftarrow o_h(1 - o_h) \sum_{k \in \text{outputs}} w_{h,k} \delta_k$$

4. Update each network weight $w_{i,j}$

$$w_{i,j} \leftarrow w_{i,j} + \Delta w_{i,j}$$

where

$$\Delta w_{i,j} = \eta \delta_j x_{i,j}$$

Aim:

Hemoglobin A1c (HbA1c) is an important measure of glucose control, which is widely applied to measure performance of diabetes care in terms of efficacy of current therapy and to make changes in that therapy. Normal HbA1c level should be lower than 7%, whereas an HbA1c higher than indicates abnormality and thus requires an adjusment in current regimen. A previous study has statistically shown that simply measuring HbA1c is associated with a lower rate of readmission in individuals with a primary diagnosis of diabetes mellitus. In light of this, the current investigation aims to build an artificial neural network for predicting the rate of readmission for a given percentage of HbA1c.

Methods

Feature Engineering

```
In [ ]: from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import log_loss
raw_df = pd.read_csv("./data/diabetes.csv")
print('Initially preprocessed Dataframe shape:', raw_df.shape)
raw_df.head()
raw_df.tail()
raw_df.info()
raw_df.isnull().sum()
#Remove duplicate encounters records for same patient, choosing the first
raw_df.drop_duplicates('patient_nbr', keep='first', inplace=True)
#Ignore the weight
raw_df = raw_df.drop(columns='weight')
#Diabetes was a diagnosis
raw_df['diag_1'] = pd.to_numeric(raw_df['diag_1'], errors='coerce')
# A1CResult were measured and Diabetes was primary diagnosis
raw_df = raw_df.dropna(axis=0, subset=['A1Cresult', 'diag_1'])
raw_df = raw_df[((raw_df['diag_1']>=250) & (raw_df['diag_1']<251))]
# Function to map A1C results to random values
def map_A1C(value):
    if value == '>8':
        return np.random.uniform(8, 10) #8
    elif value == '>7':
        return np.random.uniform(7, 8) #7
    elif value == 'Norm':
        return np.random.uniform(0, 7) #6
def map_readmission(value):
    if value == '>30':
        return 2
    elif value == '<30':
        return 1
    elif value == 'NO':
        return 0
raw_df['A1Cresult'] = raw_df['A1Cresult'].apply(map_A1C)
y = raw_df[['readmitted']].values.ravel()
cf = [col for col in raw_df.columns if raw_df[col].dtypes=='object']
le = LabelEncoder()
raw_df[cf] = raw_df[cf].apply(le.fit_transform)
X = raw_df[['diag_1', 'A1Cresult']].values
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.5, random_state=42)
# Standardize features by removing the mean and scaling to unit variance
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
mlp = MLPClassifier(random_state=0)
#Defining hyperparameters and their possible values
param_grid = {
    'hidden_layer_sizes': [(50,), (100,), (50, 50), (100,100)],
    'activation': ['tanh', 'relu', 'logistic'],
    'solver': ['sgd', 'adam'],
    'alpha': [0.001, 0.1, 1],
    'max_iter': [2000]
}
#Creating a GridSearchCV object
grid_search = GridSearchCV(mlp, param_grid, cv=5, scoring='accuracy')
```

```
#Performing the hyperparameter search
grid_search.fit(X_train_scaled, y_train)
```

Initially preprocessed Dataframe shape: (101766, 51)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 101766 entries, 0 to 101765

Data columns (total 51 columns):

#	Column	Non-Null Count	Dtype
0	id	101766 non-null	int64
1	encounter_id	101766 non-null	int64
2	patient_nbr	101766 non-null	int64
3	race	101766 non-null	object
4	gender	101766 non-null	object
5	age	101766 non-null	object
6	weight	101766 non-null	object
7	admission_type_id	101766 non-null	int64
8	discharge_disposition_id	101766 non-null	int64
9	admission_source_id	101766 non-null	int64
10	time_in_hospital	101766 non-null	int64
11	payer_code	101766 non-null	object
12	medical_specialty	101766 non-null	object
13	num_lab_procedures	101766 non-null	int64
14	num_procedures	101766 non-null	int64
15	num_medications	101766 non-null	int64
16	number_outpatient	101766 non-null	int64
17	number_emergency	101766 non-null	int64
18	number_inpatient	101766 non-null	int64
19	diag_1	101766 non-null	object
20	diag_2	101766 non-null	object
21	diag_3	101766 non-null	object
22	number_diagnoses	101766 non-null	int64
23	max_glu_serum	5346 non-null	object
24	A1Cresult	17018 non-null	object
25	metformin	101766 non-null	object
26	repaglinide	101766 non-null	object
27	nateglinide	101766 non-null	object
28	chlorpropamide	101766 non-null	object
29	glimepiride	101766 non-null	object
30	acetohexamide	101766 non-null	object
31	glipizide	101766 non-null	object
32	glyburide	101766 non-null	object
33	tolbutamide	101766 non-null	object
34	pioglitazone	101766 non-null	object
35	rosiglitazone	101766 non-null	object
36	acarbose	101766 non-null	object
37	miglitol	101766 non-null	object
38	troglitazone	101766 non-null	object
39	tolazamide	101766 non-null	object
40	examide	101766 non-null	object
41	citoglipton	101766 non-null	object
42	insulin	101766 non-null	object
43	glyburide.metformin	101766 non-null	object
44	glipizide.metformin	101766 non-null	object
45	glimepiride.pioglitazone	101766 non-null	object
46	metformin.rosiglitazone	101766 non-null	object
47	metformin.pioglitazone	101766 non-null	object
48	change	101766 non-null	object
49	diabetesMed	101766 non-null	object
50	readmitted	101766 non-null	object

dtypes: int64(14), object(37)

memory usage: 39.6+ MB

Out[]:

GridSearchCV ⓘ ?

estimator: MLPClassifier

MLPClassifier ?

Train and Test Sets

```
In [ ]: print("Best Parameters:", grid_search.best_params_)
print("Best Score:", grid_search.best_score_)
#Evaluating the best model using cross-validation: assess model's performance more robustly by training and validating it
# on different subsets of the data.
cross_val_scores = cross_val_score(grid_search.best_estimator_,X_test_scaled,y_test,cv=3,scoring='accuracy')
print("Cross-Validation Test Accuracy:", np.mean(cross_val_scores))
# Variance: Measure that helps you understand how much model's performance is spread out or varies across different subsets
# of data.
print("Variance: ",np.var(cross_val_scores))
#Calculating the Logarithmic Loss/log_loss: measures performance of classifier where input is are probabilities,
# between 0 and 1 and the output is the difference between the predicted probability distribution and the actual distribution of the outcomes.
y_pred_proba = grid_search.best_estimator_.predict_proba(X_test_scaled)
log_loss_scores = log_loss(y_test,y_pred_proba)
print("Logarithmic Loss:", log_loss_scores)
# Draw a matrix of scatter plots from the dataframe and color by y_train
X_train_dataframe = pd.DataFrame(X_train_scaled, columns=['diag_1', 'A1Cresult'])
y_train_dataframe = pd.DataFrame(y_train, columns=['readmitted'])['readmitted'].apply(map_readmission)
pd.plotting.scatter_matrix(X_train_dataframe, c=y_train_dataframe.values, figsize=(15, 15),
                           marker='o', hist_kwds={'bins': 50}, s=60, alpha=.6, cmap=mglearn.cm3)

handles = [plt.plot([],[],color=list(reversed(mglearn.cm3.colors))[i], ls="", marker=".", markersize=np.sqrt(10))[0] for i in range(3)]
labels=['>30', '<30', 'NO']
plt.legend(handles, labels, loc=(1.02,0))
plt.show()
```

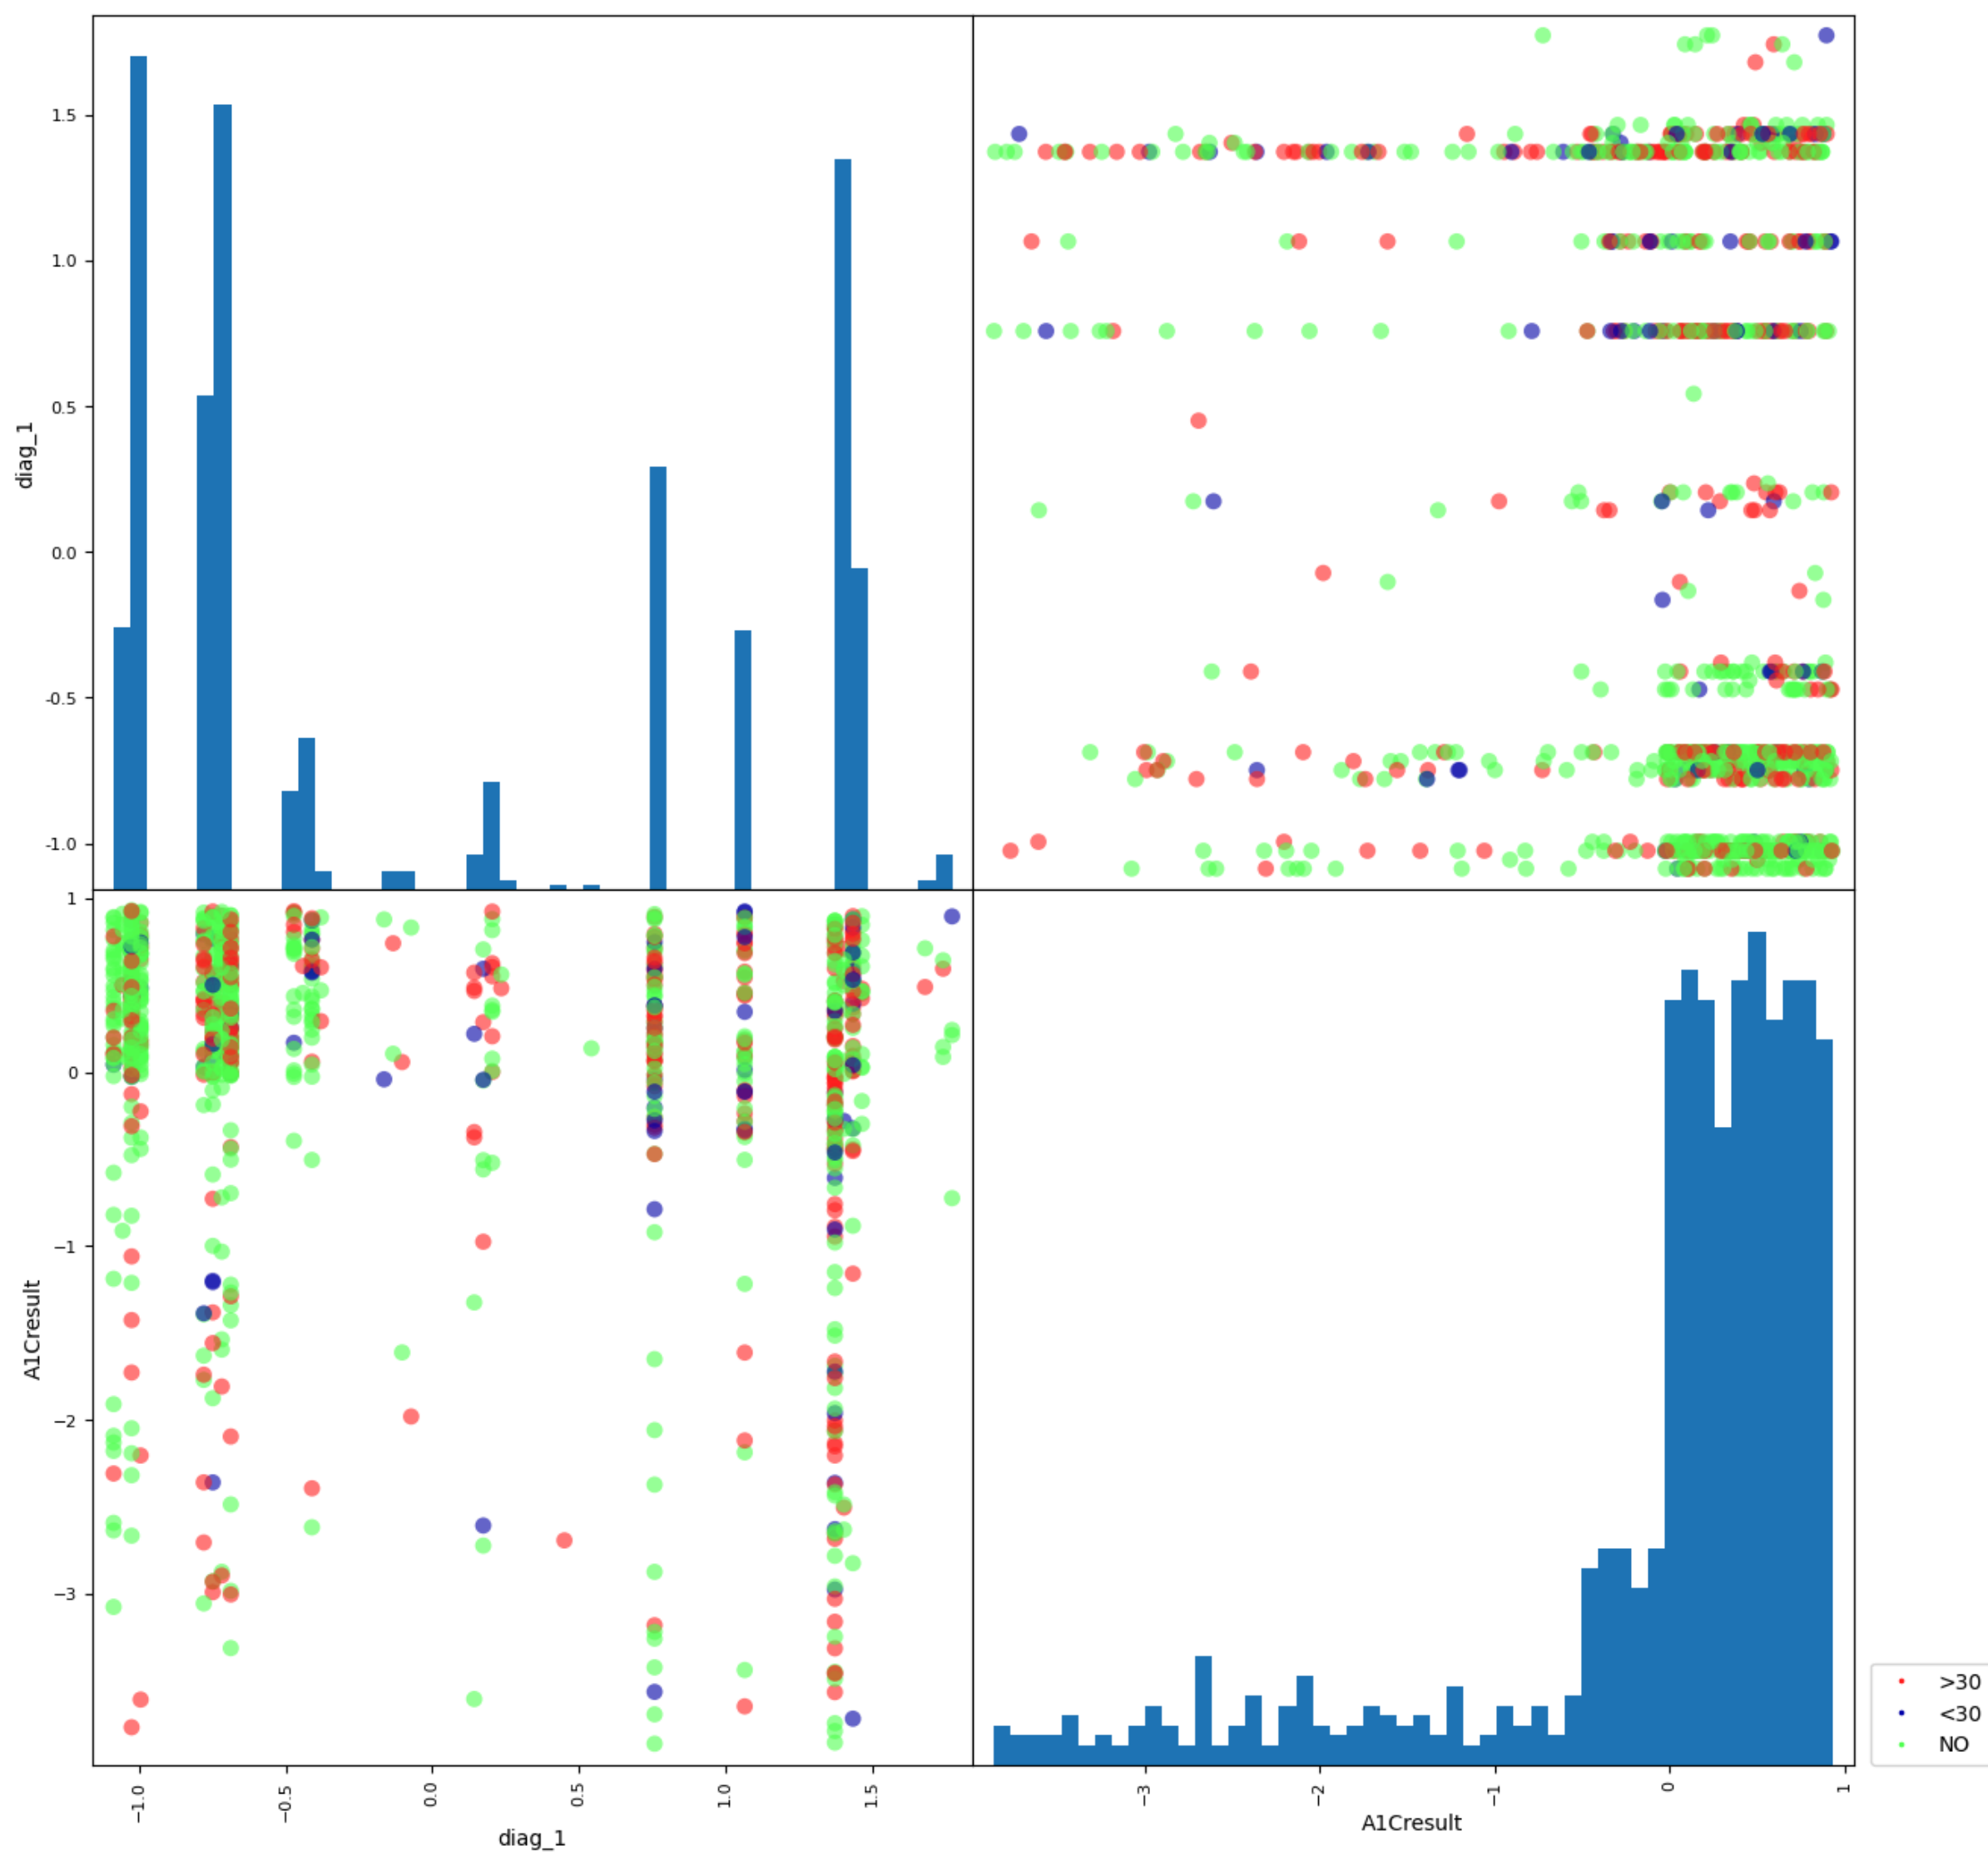
Best Parameters: {'activation': 'tanh', 'alpha': 0.001, 'hidden_layer_sizes': (50, 50), 'max_iter': 2000, 'solver': 'adam'}

Best Score: 0.6274154589371982

Cross-Validation Test Accuracy: 0.6460947753483567

Variance: 1.4319559456871993e-06

Logarithmic Loss: 0.8175268602659499

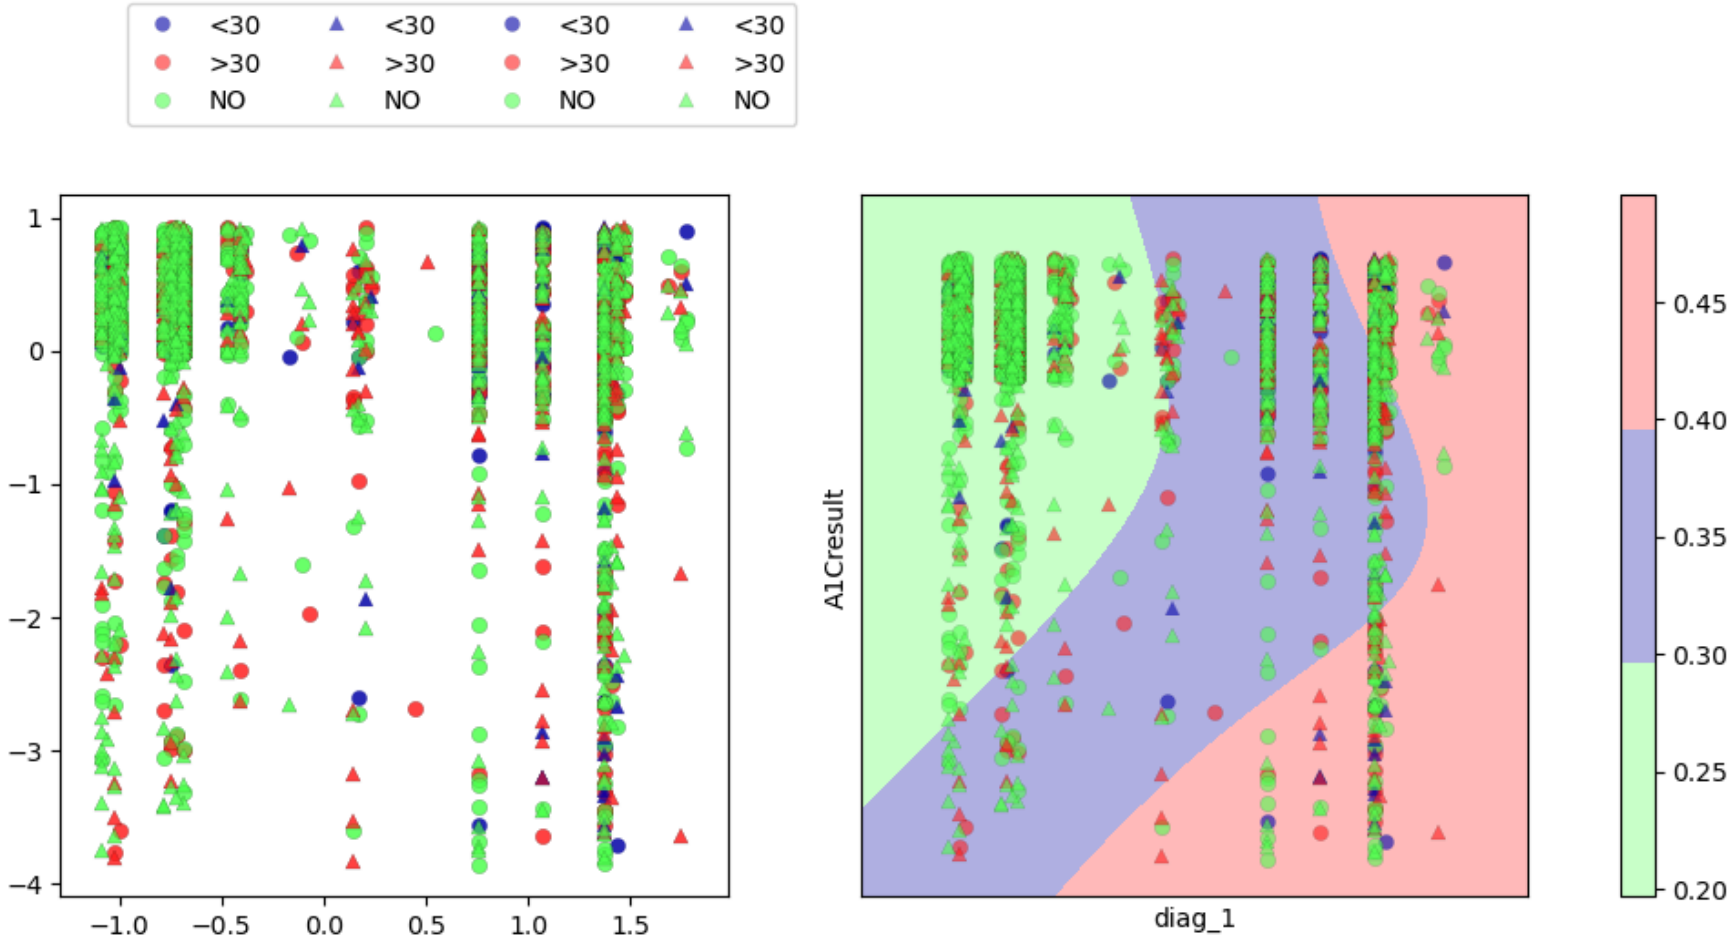


Cross validation

```
In [ ]: mlp = MLPClassifier(solver='sgd', activation=grid_search.best_params_['activation'],
                           alpha=grid_search.best_params_['alpha'], random_state=0, hidden_layer_sizes=grid_search.best_params_['hidden_layer_sizes'],
                           max_iter=grid_search.best_params_['max_iter']).fit(X_train_scaled, y_train)
# #MLP NN Learned non-linear and ragged decision boundary => model complexity high
# mglearn.plots.plot_2d_classification(mlp, scaler.fit_transform(X), fill=True, alpha=.3)
# Uncertainty estimates of predictions
fig, axes = plt.subplots(1, 2, figsize=(13, 5))
mglearn.discrete_scatter(X_train_scaled[:, 0], X_train_scaled[:, 1], y_train, markers='o', alpha=.6, s=6, markeredgewidth=.1, ax=axes[0],
                          c=[ list(reversed(mglearn.cm3.colors))[1], list(reversed(mglearn.cm3.colors))[0], list(reversed(mglearn.cm3.colors))[2]])
mglearn.discrete_scatter(X_test_scaled[:, 0], X_test_scaled[:, 1], y_test, markers='^', alpha=.6, s=6, markeredgewidth=.1, ax=axes[0],
                          c=[ list(reversed(mglearn.cm3.colors))[1], list(reversed(mglearn.cm3.colors))[0], list(reversed(mglearn.cm3.colors))[2]])
scores_image = mglearn.tools.plot_2d_scores(mlp, scaler.fit_transform(X), ax=axes[1], alpha=.3, cm=mglearn.cm3, function='predict_proba')
for ax in axes:
    # plot training and test points
    mglearn.discrete_scatter(X_train_scaled[:, 0], X_train_scaled[:, 1], y_train, markers='o', alpha=.6, s=6, markeredgewidth=.1, ax=ax,
                              c=[ list(reversed(mglearn.cm3.colors))[1], list(reversed(mglearn.cm3.colors))[0], '#50ff50'])
    mglearn.discrete_scatter(X_test_scaled[:, 0], X_test_scaled[:, 1], y_test, markers='^', alpha=.6, s=6, markeredgewidth=.1, ax=ax,
                              c=[ list(reversed(mglearn.cm3.colors))[1], list(reversed(mglearn.cm3.colors))[0], list(reversed(mglearn.cm3.colors))[2]])

    plt.xlabel("diag_1")
    plt.ylabel("A1Cresult")
cbar = plt.colorbar(scores_image, ax=axes.tolist())
axes[0].legend(ncol=4, loc=(.1, 1.1))
# # #Accuracy of the MLP is quite good, but not as good as the other models. Likely due to scaling of the data.
# # #NN expect all input features to vary in a similar way, and ideally to have a mean=0, and variance=1.
print("Accuracy on training set: {:.2f}".format(mlp.score(X_train_scaled, y_train)))
print("Accuracy on test set: {:.2f}".format(mlp.score(X_test_scaled, y_test)))
```

Accuracy on training set: 0.63
Accuracy on test set: 0.65



Conclusion

- ANN trained
- Best Parameters: {'activation': 'tanh', 'alpha': 0.001, 'hidden_layer_sizes': (50, 50), 'max_iter': 2000, 'solver': 'sgd'}

- Best Score: 0.6283816425120774
- Cross-Validation Test Accuracy: 0.65
- - ANN predicted correctly 65% of the time.
- Variance: 2.32e-07
- - Low variance suggest good consitent and prediction accross training and test set
- Logarithmic Loss: 0.81, indicates that probabilities are not very aligned with actual outcomes. Lower log loss values are better, with 0 indicating a perfect model. 0.81 suggests a considerable level of uncertainty in the predictions, as supported by the predicted probabilty map.